

SPIMAP: Species Informed Maximum A Posteriori Gene Tree Reconstruction

Documentation for the SPIMAP software package

April 8, 2011

Author: Matthew D. Rasmussen (rasmus@mit.edu)

Software website: <http://compbio.mit.edu/spimap>

citation:

Matthew D. Rasmussen and Manolis Kellis. *A Bayesian Approach for Fast and Accurate Gene Tree Reconstruction*. Molecular Biology and Evolution. 2010. doi: 10.1093/molbev/msq189

Contents

| | | |
|----------|---|----------|
| 1 | File formats | 3 |
| 1.1 | Sequence alignment format (*.align) | 3 |
| 1.2 | Species tree file format (*.stree) | 3 |
| 1.3 | Gene to species name mapping file format (*.smap) | 4 |
| 1.4 | Reconciliation file format (*.recon) | 5 |
| 1.5 | SPIMAP model parameters file format (*.params) | 6 |

```

>KLLA0C08239g
ATGAGTCTCAAACGTGTAGTTGTCAGTGGTCTTGGGGCCTACACGCCCTTGGTTCTACAGTTTCAAAGTCTTGGGCAGG
TTTGCTT—GCTGCTAAGCAATCACTAATACCCTTAGATGCTTTCTACAACAGAGAA—GACTTTGCAAAAGTGAAAA
AGTTGGTCCCAGTACAGCAGTGAAGTAGGTTACAT—
>ADL072C
—ATGCATCCCCGAGTGGTCGTGACCGGCATTGGGTGCTATACTCCTCTGGGGCCGTCGCTAGCCCAGTCTTGAAGGA
GCTGTTG—CGAGGGACGAGCGGCCTTGTGAGGCTGCAAGATCTGGCAGAGTACGAGGGCGATTACAAACCACTGTGCA
GGCTTATATCCGGTGATCTTCGAGTCGGGAAAGTTGGATTGAG—
>kwal_5828
—ATGACTTCCAGAGTCGTTGTTACTGGGCTTGGTGCTATCACTCCACTTGGGAGGACTGTTTCCGAGTCATGGAGAGC
TTTATTG—GCAGGCAAGTCCGGAATTCGTCCCATTCGCGATCTTCCC—AATGCTAAAAGCTACGAAG
GACACTGTCTGCATCTGTTGCCGTTGCAGACATTCCTGATTTG—GATCCA—

```

Figure 1: **Example *.align file.** Three gene DNA sequences are given, each with 240 sites.

1 File formats

1.1 Sequence alignment format (*.align)

SPIMAP uses the FASTA file format (http://en.wikipedia.org/wiki/FASTA_format) for sequences alignments. The file extension is not important and many different extensions are in common use (*.fa, *.mfa, *.fasta, *.align).

Each line starting with ">" is used to indicate a gene name (Figure 1). The entire line after the ">" is used as the gene name. The aligned sequence of is then given on the following lines. It may be wrapped to any number of columns (no not wrapped at all). Gaps in the alignment are represented with the "-" character.

At this time (version 1.1), SPIMAP can only use DNA sequences. The sequence can be in both upper case and lower case (SPIMAP ignores case) and degeneracy codes can be used ("NnRrYyWwSsKkMmBbD-dHhVv"), however at this time SPIMAP treats all degeneracies as completely missing data ("N").

1.2 Species tree file format (*.stree)

Species trees should be specified using the Newick file format. See http://en.wikipedia.org/wiki/Newick_format for details. Beyond the newick format, SPIMAP has only a few additional requirements. First, the species names given in the species tree should match those given in the SMAP file (Section 1.3). Second, the branch lengths of the species tree should be expressed in units of time (Figure 2). Any unit of time can be used (e.g. millions of years, generations, relative units, etc). The only requirement is that the duplication and loss rates are also expressed in compatible units. Therefore, if branch lengths are in *millions of years*, the duplication rate (specified by spimap's "-D" option) should be in units of duplications/gene/*million years*.




Figure 2: **Example *.stree file and corresponding tree.** This file (left) specifies the species tree (right) using the newick file format. Branch lengths should be expressed in units of time (e.g. millions of years). Ancestral nodes can also optionally be named (the names “n1”, “n2”, etc are used in this example).

Naming ancestral nodes. SPIMAP also supports naming ancestral nodes in the species tree using the newick format. For example, the parental node of `human` and `chimp` can be named `primate` using the following syntax:

```
((human:5,chimp:5)primate:70,mouse:75)mammal;
```

If ancestral nodes are named, they will be used in the output of the reconciliation mapping (Section 1.4).

1.3 Gene to species name mapping file format (*.smap)

SPIMAP uses a special file format (*.smap) to specify which genes belong to which species. Each line contains two tab-delimited fields:

1. pattern matching a gene ID
2. species ID

Only 3 types of gene ID patterns are supported. The pattern can either be an exact matching string, a prefix (denoted "text*"), or a suffix (denoted "*text"). The "*" is the only special wildcard character.

| | |
|--------|------|
| A* | agos |
| orf19* | calb |
| CDUG_* | cdub |
| CAGL* | cgl |
| IPF_* | cgl |
| CGUG_* | cgui |
| sbay_* | sbay |
| scas_* | scas |
| smik_* | smik |
| spar_* | spar |
| SCP* | spom |
| YALI* | ylip |
| Y* | scer |
| Q* | scer |

Figure 3: **Example *.smap file.** This file specifies how to map gene names to their corresponding species. The first column indicates a gene name pattern (in this case a prefix) and the second column specifies a species name. Note: this example only gives a partial list of the species in Figure 2.

The species ID should be the same as the ones used in the species tree files. All patterns and IDs are case-sensitive.

When mapping a gene name to a species name all exact matches are processed first. If no exact match is found, the patterns are then processed in the same order as they appear in the file until a match is found. For example in the SMAP file given in Figure 4, the gene ID "YALI123" should match the species "ylip", instead of "yeast", because the pattern "YALI*" occurs before "Y*".

1.4 Reconciliation file format (*.recon)

When SPIMAP's "-r" option is used, the reconciliation found for the gene tree and species is saved to a file "OUTPUT_PREFIX.recon". The reconciliation file format is tab-delimited, where each line has three fields:

1. gene node ID.
2. species node ID.
3. event (one of the following: "gene", "spec", "dup")

Each line specifies the mapping of one node in the gene tree (field 1) to one node or branch in the species tree (field 2). Branches are indicated using the node ID directly below it (i.e. the younger of the two incident nodes). The lines can be given in any order.

If the gene node is a leaf, it will map to a leaf in the species tree and the event field will contain the event "gene". All internal nodes of the gene tree are marked either as speciations (event "spec") or duplications (event "dup"). Speciation nodes map directly to the indicated species node, and duplication nodes map to the indicated species branch. The time of the duplication along the species branch is not indicated in this file format nor is it inferred by SPIMAP.

| | | |
|--------------|------|------|
| KLLA0C08239g | klac | gene |
| ADL072C | agos | gene |
| kwal_5828 | kwal | gene |
| CAGL0J02970g | cgl | gene |
| scas_g715.48 | scas | gene |
| smik_6662 | smik | gene |
| sbay_7039 | sbay | gene |
| smik_6659 | smik | gene |
| sbay_7037 | sbay | gene |
| YER061C | scer | gene |
| spar_6281 | spar | gene |
| n10 | n5 | spec |
| n9 | n7 | spec |
| n8 | n6 | spec |
| n7 | n5 | spec |
| n6 | n5 | dup |
| n5 | n3 | spec |
| n4 | n3 | spec |
| n3 | n9 | spec |
| n2 | n8 | spec |
| n1 | n2 | spec |




Figure 4: **Example *.recon file.** The reconciliation file format (left) specifies how all the nodes in a gene tree (right) map to the nodes and branches in the species tree (see Figure 2). Notice that gene node “n6” (red dot) represents a duplication event along species branch “n5” (shown Figure 2). The gene tree and species tree have their own name space (“n5” in the gene tree is not the same as “n5” in the species tree).

If gene IDs are not given to the ancestral nodes of a gene tree or species tree, SPIMAP will by default name them with “nXXX” where XXX is the preorder traversal of the internal nodes.

1.5 SPIMAP model parameters file format (*.params)

SPIMAP has several parameters for its substitution rates model. These parameters are learned by the `spimap-train-rates` program, which saves the parameters in a custom *.params file format. The `spimap` program reads these parameters using the “-p” option. Most uses of SPIMAP do not require understanding the contents of a *.params file.

The *.params file format is tab-delimited and each line is processed one at a time.

If the first field of a line is the word “baserate”, then the remaining two fields are interpreted as floating point values α_G and β_G , which are the two parameters, shape and scale, of the inverse-gamma distributed gene-specific rate.

If the first field of the line does not match “baserate”, then the first field indicates a species tree branch and the remaining two fields are interpreted as floating point values α_i and β_i , which are the two parameters, shape and scale, of the gamma distributed species-specific rate. Each branch is indicated by its more recent

| | | |
|----------|----------------|---------------|
| baserate | 6.98457288742 | 5.98457288742 |
| 1 | 3.28887700831 | 394.209221588 |
| 2 | 4.64684152603 | 551.109741211 |
| 3 | 1.13027572632 | 164.191940308 |
| 4 | 0.610769152641 | 75.0393371582 |
| 5 | 7.14405012131 | 927.631103516 |
| 6 | 2.96983885765 | 238.195861816 |
| 7 | 5.63683271408 | 632.264831543 |
| 8 | 0.974860072136 | 94.9837493896 |
| 9 | 0.856632292271 | 78.6899032593 |
| 10 | 4.64683914185 | 544.528686523 |
| 11 | 1.92581880093 | 271.891052246 |
| 12 | 3.84569692612 | 624.703308105 |
| 13 | 3.14617466927 | 335.446655273 |
| 14 | 0.699178874493 | 84.1814575195 |
| 15 | 0.746283352375 | 137.345901489 |
| scer | 8.42576217651 | 763.305847168 |
| ctro | 6.70220327377 | 999.845153809 |
| scas | 9.14448356628 | 1253.45031738 |
| agos | 8.84074497223 | 801.648925781 |
| sbay | 6.95680332184 | 1048.7590332 |
| kwal | 14.3321857452 | 1962.9083252 |
| dhan | 15.7483224869 | 2699.00878906 |
| smik | 10.2562847137 | 1143.78076172 |
| cgl | 9.81903266907 | 1015.43951416 |
| spar | 5.80616807938 | 799.18963623 |
| calb | 8.38038921356 | 1233.68322754 |
| lelo | 9.40990924835 | 973.772583008 |
| cpar | 9.43262672424 | 1184.28100586 |
| klac | 6.6709280014 | 767.418823242 |
| clus | 8.37989234924 | 881.762878418 |
| cgui | 11.9692239761 | 1187.47314453 |

Figure 5: **Example *.params file.** The *.params file contains the parameters for SPIMAP's substitution rate model.

node. Ancestral nodes are indicated by an integer, where are assigned in pretraversal order.